


ASTRA-SIM Description



Saeed Rashidi

Ph.D. Student, School of Electrical & Computer Engineering
Georgia Institute of Technology

saeed.rashidi@gatech.edu


Acknowledgments: Srinivas Sridharan (Meta), Sudarshan Srinivasan (Intel)

Overview

How to Model and Evaluate the Communication Effect

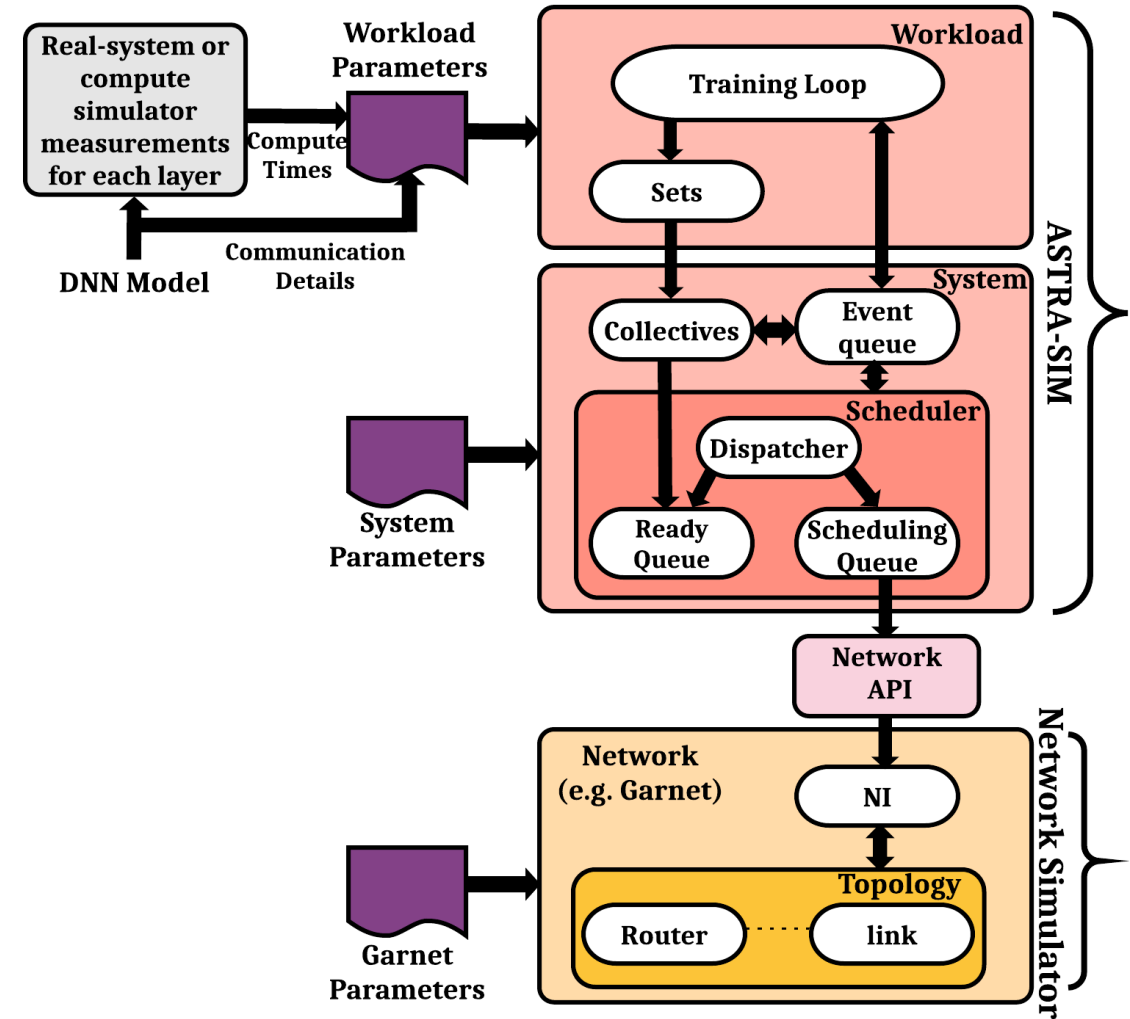
- It is a complex problem and can be viewed as three layers :

- 1. Workload layer (the training loop):
 - Parallelism approach
 - Compute power
 - Communication size & type and dependency order
- 2. System layer:
 - Collective communication algorithm
 - Chunk size, schedule of collectives
- 3. Network layer:
 - Physical topology
 - Congestion control, communication protocol
 - Link BW, latency, buffers, routing algorithm




ASTRA-SIM Architecture

- **Workload layer:**
 - Supports Data-Parallel, Model-Parallel, Hybrid-Parallel
 - Easy to add new arbitrary training loop
 - **Two** workload front-ends:
 - Text-based workload engine (vanilla)
 - Execution graph engine (**New**)
- **System:**
 - Ring based, Tree-based, Direct based, and multi-processor
 - Easy to add new collective communication
- **Network:**
 - Analytical:
 - Supports hierarchical topologies
 - Each level in hierarchy can be switch, ring, FC...
 - <https://github.com/astra-sim/analytical/tree/develop>
 - GARNET:
 - Supports switch-based and torus-based topologies
 - Supports credit-based flow control
 - https://github.com/georgia-tech-synergy-lab/gem5_astra/tree/main
 - NS3 (**New**):
 - Supports datacenter network networks (e.g., RoCE)
 - <https://github.com/DartingMelody/ns3-interface>



ASTRA-SIM Runtime Structure

- Each NPU is represented through separate instance of Workload, System, and Network API.
- Network API class is implemented by the network backend.



ASTRA-SIM Directory

rashidi1saeed Merge pull request #44 from srinivas212/master ...			✖ f7e54a8 15 hours ago	🕒 335 commits
📁 .github/workflows	Added GitHub Actions (#30)			8 months ago
📁 astra-sim	CSV Writer Updated			19 hours ago
📁 build	gem5 is now compatible with new changes			18 hours ago
📁 docs/images	updated			9 months ago
📁 examples	run_multi example script bug fixed			18 hours ago
📁 extern	Merge pull request #43 from astra-sim/saeed_astra_dev			15 hours ago
📁 inputs	Update README.md			17 hours ago
📁 scripts/workload_generator	-.TESTED:-			last month
📁 test	Fix formatting using clang-format			9 months ago
📄 .clang-format	Use PyTorch .clang-format			9 months ago
📄 .clang-tidy	Added GitHub Actions (#30)			8 months ago
📄 .gitignore	Added GitHub Actions (#30)			8 months ago
📄 .gitmodules	Add scale sim v2 submodule			last month
📄 CMakeLists.txt	CMAKELists updated			yesterday
📄 CODEOWNERS	Modify CODEOWNERS			16 hours ago
📄 LICENSE	Update LICENSE			9 months ago
📄 README.md	Update README.md			9 months ago


Network Input

Garnet network input

```
sample_torus.txt
1 num-npus: 12
2 num-packages: 6
3 package-rows: 3
4 topology: Torus3D
5 local-rings: 2
6 vertical-rings: 1
7 horizontal-rings: 1
8 flit-width: 2048
9 local-packet-size: 4096
10 package-packet-size: 4096
11 tile-link-width: 256
12 package-link-width: 256
13 vcs-per-vnet: 50
14 routing-algorithm: Ring_XY
15 router-latency: 1
16 local-link-latency: 90
17 package-link-latency: 200
18 buffers-per-vc: 5000
19 local-link-efficiency: 1.0
20 package-link-efficiency: 1.0
21
```

Analytical network input

```
sample_Torus3D.json
1 {
2   "topology-name": "Hierarchical",
3   "topologies-per-dim": ["Ring", "Ring", "Ring"],
4   "dimension-type": ["N", "N", "N"],
5   "dimensions-count": 3,
6   "units-count": [2, 2, 3],
7   "links-count": [2, 2, 2],
8   "link-latency": [10, 100, 100],
9   "link-bandwidth": [32, 16, 16],
10  "nic-latency": [0, 0, 0],
11  "router-latency": [0, 0, 0],
12  "hbm-latency": [500, 500, 500],
13  "hbm-bandwidth": [370, 370, 370],
14  "hbm-scale": [0, 0, 0]
15 }
```



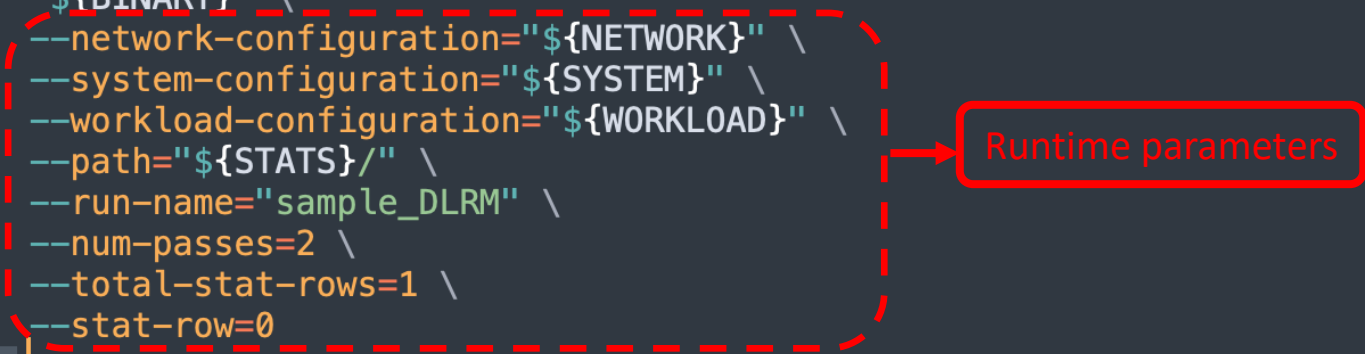
ASTRA-SIM Run Script

A Sample Run Script

```
run_DLRLM_analytical.sh x
1  #! /bin/bash -v
2
3  # Absolute path to this script
4  SCRIPT_DIR=$(dirname "$(realpath $0)")
5
6  # Absolute paths to useful directories
7  BINARY="${SCRIPT_DIR:?}"/../build/astra_analytical/build/AnalyticalAstra/bin/AnalyticalAstra
8  NETWORK="${SCRIPT_DIR:?}"/../inputs/network/analytical/sample_Torus3D.json
9  SYSTEM="${SCRIPT_DIR:?}"/../inputs/system/sample_torus_sys.txt
10 WORKLOAD="${SCRIPT_DIR:?}"/../inputs/workload/DLRLM_HybridParallel.txt
11 STATS="${SCRIPT_DIR:?}"/results/run_DLRLM_analytical
12
13 rm -rf "${STATS}"
14 mkdir "${STATS}"
15
16 "${BINARY}" \
17 --network-configuration="${NETWORK}" \
18 --system-configuration="${SYSTEM}" \
19 --workload-configuration="${WORKLOAD}" \
20 --path="${STATS}/" \
21 --run-name="sample_DLRLM" \
22 --num-passes=2 \
23 --total-stat-rows=1 \
24 --stat-row=0
25
26
```

A Sample Run Script

```
run_DLRM_analytical.sh x
1  #!/bin/bash -v
2
3  # Absolute path to this script
4  SCRIPT_DIR=$(dirname "$(realpath $0)")
5
6  # Absolute paths to useful directories
7  BINARY="${SCRIPT_DIR:?}"/../build/astra_analytical/build/AnalyticalAstra/bin/AnalyticalAstra
8  NETWORK="${SCRIPT_DIR:?}"/../inputs/network/analytical/sample_Torus3D.json
9  SYSTEM="${SCRIPT_DIR:?}"/../inputs/system/sample_torus_sys.txt
10 WORKLOAD="${SCRIPT_DIR:?}"/../inputs/workload/DLRM_HybridParallel.txt
11 STATS="${SCRIPT_DIR:?}"/results/run_DLRM_analytical
12
13 rm -rf "${STATS}"
14 mkdir "${STATS}"
15
16 "${BINARY}" \
17 --network-configuration="${NETWORK}" \
18 --system-configuration="${SYSTEM}" \
19 --workload-configuration="${WORKLOAD}" \
20 --path="${STATS}/" \
21 --run-name="sample_DLRM" \
22 --num-passes=2 \
23 --total-stat-rows=1 \
24 --stat-row=0
25
26
```



ASTRA-SIM Reports

Overall Results

- Endtoend.csv.

Layer name

	Layer Name	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	
1				fwd compute	wg compute	ig compute	fwd exposed	wg exposed	cig exposed	cc	fwd total com	wg total com	ig total com	n	workload fini	total comp	total exposed comm
2	conv1		sample_Resnet	26.006	64.582	0	0	17.364	0	0	17.366	0	4875.201	2164.9	2710.301		
3	layer_64_1_c		sample_Resnet	6.912	14.976	7.296	0	0	0	0	4.796	0	4875.201				
4	layer_64_1_c		sample_Resnet	6.912	14.976	6.912	0	0	0	0	27.08	0	4875.201				
5	layer_64_1_c		sample_Resnet	21.888	28.288	20.736	0	0	0	0	47.906	0	4875.201				
6	layer_64_1_c		sample_Resnet	6.912	14.976	7.296	0	0	0	0	13.648	0	4875.201				
7	layer_64_2_c		sample_Resnet	7.296	19.968	6.912	0	0	0	0	10.166	0	4875.201				
8	layer_64_2_c		sample_Resnet	21.888	28.288	20.736	0	0	0	0	20.102	0	4875.201				
9	layer_64_2_c		sample_Resnet	6.912	14.976	7.296	0	0	0	0	22.048	0	4875.201				
10	layer_64_3_c		sample_Resnet	7.296	19.968	6.912	0	0	0	0	30.082	0	4875.201				
11	layer_64_3_c		sample_Resnet	21.888	28.288	20.736	0	0	0	0	11.334	0	4875.201				
12	layer_64_3_c		sample_Resnet	6.912	14.976	7.296	0	0	0	0	7.526	0	4875.201				
13	layer_128_1		sample_Resnet	5.184	12.288	5.184	0	0	0	0	36.03	0	4875.201				
14	layer_128_1		sample_Resnet	7.296	19.968	7.04	0	0	0	0	9.08	0	4875.201				
15	layer_128_1		sample_Resnet	12.96	13.312	11.68	0	0	0	0	28.13	0	4875.201				
16	layer_128_1		sample_Resnet	4.672	10.24	5.184	0	0	0	0	13.272	0	4875.201				
17	layer_128_2		sample_Resnet	5.184	8.192	4.672	0	0	0	0	30.868	0	4875.201				
18	layer_128_2		sample_Resnet	12.96	13.312	11.68	0	0	0	0	33.952	0	4875.201				
19	layer_128_2		sample_Resnet	4.672	10.24	5.184	0	0	0	0	101.056	0	4875.201				

Overall Results

- Endtoend.csv.

Run name

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	
1			fwd compute	wg compute	ig compute	fwd exposed	wg exposed	cig exposed	cc	fwd total com	wg total com	ig total com	n	workload fini	total comp	total exposed comm
2	conv1	sample_Resnet	26.006	64.582	0	0	17.364	0	0	17.366	0	4875.201	2164.9	2710.301		
3	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	4.796	0	4875.201				
4	layer_64_1_c	sample_Resnet	6.912	14.976	6.912	0	0	0	0	27.08	0	4875.201				
5	layer_64_1_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	47.906	0	4875.201				
6	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	13.648	0	4875.201				
7	layer_64_2_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	10.166	0	4875.201				
8	layer_64_2_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	20.102	0	4875.201				
9	layer_64_2_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	22.048	0	4875.201				
10	layer_64_3_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	30.082	0	4875.201				
11	layer_64_3_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	11.334	0	4875.201				
12	layer_64_3_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	7.526	0	4875.201				
13	layer_128_1	sample_Resnet	5.184	12.288	5.184	0	0	0	0	36.03	0	4875.201				
14	layer_128_1	sample_Resnet	7.296	19.968	7.04	0	0	0	0	9.08	0	4875.201				
15	layer_128_1	sample_Resnet	12.96	13.312	11.68	0	0	0	0	28.13	0	4875.201				
16	layer_128_1	sample_Resnet	4.672	10.24	5.184	0	0	0	0	13.272	0	4875.201				
17	layer_128_2	sample_Resnet	5.184	8.192	4.672	0	0	0	0	30.868	0	4875.201				
18	layer_128_2	sample_Resnet	12.96	13.312	11.68	0	0	0	0	33.952	0	4875.201				
19	layer_128_2	sample_Resnet	4.672	10.24	5.184	0	0	0	0	101.056	0	4875.201				

Overall Results

- Endtoend.cs Compute times (us)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	
1			fwd compute	wg compute	ig compute	fwd exposed	wg exposed	ig exposed	cc	fwd total com	wg total com	ig total com	n	workload fini	total comp	total exposed comm
2	conv1	sample_Resnet	26.006	64.582	0	0	17.364	0	0	0	17.366	0	4875.201	2164.9	2710.301	
3	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	4.796	0	4875.201			
4	layer_64_1_c	sample_Resnet	6.912	14.976	6.912	0	0	0	0	0	27.08	0	4875.201			
5	layer_64_1_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	47.906	0	4875.201			
6	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	13.648	0	4875.201			
7	layer_64_2_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	0	10.166	0	4875.201			
8	layer_64_2_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	20.102	0	4875.201			
9	layer_64_2_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	22.048	0	4875.201			
10	layer_64_3_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	0	30.082	0	4875.201			
11	layer_64_3_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	11.334	0	4875.201			
12	layer_64_3_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	7.526	0	4875.201			
13	layer_128_1	sample_Resnet	5.184	12.288	5.184	0	0	0	0	0	36.03	0	4875.201			
14	layer_128_1	sample_Resnet	7.296	19.968	7.04	0	0	0	0	0	9.08	0	4875.201			
15	layer_128_1	sample_Resnet	12.96	13.312	11.68	0	0	0	0	0	28.13	0	4875.201			
16	layer_128_1	sample_Resnet	4.672	10.24	5.184	0	0	0	0	0	13.272	0	4875.201			
17	layer_128_2	sample_Resnet	5.184	8.192	4.672	0	0	0	0	0	30.868	0	4875.201			
18	layer_128_2	sample_Resnet	12.96	13.312	11.68	0	0	0	0	0	33.952	0	4875.201			
19	layer_128_2	sample_Resnet	4.672	10.24	5.184	0	0	0	0	0	101.056	0	4875.201			

Overall Results

- Endtoend.csv.

Raw communication times (us)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1			fwd compute	wg compute	ig compute	fwd exposed	wg exposed	ig exposed	cc	fwd total com	wg total com	ig total com	workload fini	total comp	total exposed comm
2	conv1	sample_Resnet	26.006	64.582	0	0	17.364	0	0	0	17.366	0	4875.201	2164.9	2710.301
3	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	4.796	0	4875.201		
4	layer_64_1_c	sample_Resnet	6.912	14.976	6.912	0	0	0	0	0	27.08	0	4875.201		
5	layer_64_1_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	47.906	0	4875.201		
6	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	13.648	0	4875.201		
7	layer_64_2_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	0	10.166	0	4875.201		
8	layer_64_2_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	20.102	0	4875.201		
9	layer_64_2_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	22.048	0	4875.201		
10	layer_64_3_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	0	30.082	0	4875.201		
11	layer_64_3_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	11.334	0	4875.201		
12	layer_64_3_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	7.526	0	4875.201		
13	layer_128_1	sample_Resnet	5.184	12.288	5.184	0	0	0	0	0	36.03	0	4875.201		
14	layer_128_1	sample_Resnet	7.296	19.968	7.04	0	0	0	0	0	9.08	0	4875.201		
15	layer_128_1	sample_Resnet	12.96	13.312	11.68	0	0	0	0	0	28.13	0	4875.201		
16	layer_128_1	sample_Resnet	4.672	10.24	5.184	0	0	0	0	0	13.272	0	4875.201		
17	layer_128_2	sample_Resnet	5.184	8.192	4.672	0	0	0	0	0	30.868	0	4875.201		
18	layer_128_2	sample_Resnet	12.96	13.312	11.68	0	0	0	0	0	33.952	0	4875.201		
19	layer_128_2	sample_Resnet	4.672	10.24	5.184	0	0	0	0	0	101.056	0	4875.201		

Overall Results

- Endtoend.csv.

Exposed communication times (us)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1			fwd compute	wg compute	ig compute	fwd exposed	wg exposed	ig exposed	cc	fwd total com	wg total com	ig total com	workload fini	total comp	total exposed comm
2	conv1	sample_Resnet	26.006	64.582	0	0	17.364	0	0	0	17.366	0	4875.201	2164.9	2710.301
3	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	4.796	0	4875.201		
4	layer_64_1_c	sample_Resnet	6.912	14.976	6.912	0	0	0	0	0	27.08	0	4875.201		
5	layer_64_1_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	47.906	0	4875.201		
6	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	13.648	0	4875.201		
7	layer_64_2_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	0	10.166	0	4875.201		
8	layer_64_2_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	20.102	0	4875.201		
9	layer_64_2_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	22.048	0	4875.201		
10	layer_64_3_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	0	30.082	0	4875.201		
11	layer_64_3_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	0	11.334	0	4875.201		
12	layer_64_3_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	0	7.526	0	4875.201		
13	layer_128_1	sample_Resnet	5.184	12.288	5.184	0	0	0	0	0	36.03	0	4875.201		
14	layer_128_1	sample_Resnet	7.296	19.968	7.04	0	0	0	0	0	9.08	0	4875.201		
15	layer_128_1	sample_Resnet	12.96	13.312	11.68	0	0	0	0	0	28.13	0	4875.201		
16	layer_128_1	sample_Resnet	4.672	10.24	5.184	0	0	0	0	0	13.272	0	4875.201		
17	layer_128_2	sample_Resnet	5.184	8.192	4.672	0	0	0	0	0	30.868	0	4875.201		
18	layer_128_2	sample_Resnet	12.96	13.312	11.68	0	0	0	0	0	33.952	0	4875.201		
19	layer_128_2	sample_Resnet	4.672	10.24	5.184	0	0	0	0	0	101.056	0	4875.201		

Overall Results

- Endtoend.csv.

Total compute & exposed communication times across all layers (us)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1			fwd compute	wg compute	ig compute	fwd exposed	wg exposed	ig exposed	cc	fwd total com	wg total com	ig total com	workload fini	total comp	total exposed comm
2	conv1	sample_Resnet	26.006	64.582	0	0	17.364	0	0	17.366	0	4875.201		2164.9	2710.301
3	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	4.796	0	4875.201			
4	layer_64_1_c	sample_Resnet	6.912	14.976	6.912	0	0	0	0	27.08	0	4875.201			
5	layer_64_1_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	47.906	0	4875.201			
6	layer_64_1_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	13.648	0	4875.201			
7	layer_64_2_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	10.166	0	4875.201			
8	layer_64_2_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	20.102	0	4875.201			
9	layer_64_2_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	22.048	0	4875.201			
10	layer_64_3_c	sample_Resnet	7.296	19.968	6.912	0	0	0	0	30.082	0	4875.201			
11	layer_64_3_c	sample_Resnet	21.888	28.288	20.736	0	0	0	0	11.334	0	4875.201			
12	layer_64_3_c	sample_Resnet	6.912	14.976	7.296	0	0	0	0	7.526	0	4875.201			
13	layer_128_1	sample_Resnet	5.184	12.288	5.184	0	0	0	0	36.03	0	4875.201			
14	layer_128_1	sample_Resnet	7.296	19.968	7.04	0	0	0	0	9.08	0	4875.201			
15	layer_128_1	sample_Resnet	12.96	13.312	11.68	0	0	0	0	28.13	0	4875.201			
16	layer_128_1	sample_Resnet	4.672	10.24	5.184	0	0	0	0	13.272	0	4875.201			
17	layer_128_2	sample_Resnet	5.184	8.192	4.672	0	0	0	0	30.868	0	4875.201			
18	layer_128_2	sample_Resnet	12.96	13.312	11.68	0	0	0	0	33.952	0	4875.201			
19	layer_128_2	sample_Resnet	4.672	10.24	5.184	0	0	0	0	101.056	0	4875.201			

Overall Results

- Detailed.csv

Average chunk queueing delay per each collective phase (us)

	A	B	C	D	E	F	G	H	I	J	K	L	M
1			queuing d	queuing delay phase 1	queuing delay phase 2	queuing delay phase 3	queuing delay phase 4	queuing delay phase 5	network delay phase 1	network delay phase 2	network delay phase 3	network delay phase 4	network delay phase 5
2	conv1	sample_Resnet	0	1.8795	0.5455	0	0.137	2.3575	0.035	0.018	0.005	0.018	0.035
3	layer_64_1	sample_Resnet	0	0.453	0.03925	0	0.02725	0.2425	0.06	0.03	0.008	0.03	0.06
4	layer_64_1	sample_Resnet	0	2.9225	2.3655	0	0.065	7.1915	0.015	0.008	0.002	0.008	0.015
5	layer_64_1	sample_Resnet	0	6.8005	2.5125	1.2645	3.2625	6.7425	0.135	0.068	0.017	0.068	0.135
6	layer_64_1	sample_Resnet	0	2.013	0.412	0	0.227	2.536	0.06	0.03	0.008	0.03	0.06
7	layer_64_2	sample_Resnet	0	2.889	0.03925	0	0.02725	0.2425	0.06	0.03	0.008	0.03	0.06
8	layer_64_2	sample_Resnet	0	5.9785	0.08275	0	0.07075	0.5485	0.135	0.068	0.017	0.068	0.135
9	layer_64_2	sample_Resnet	0	1.901	2.947	0	0.227	4.313	0.06	0.03	0.008	0.03	0.06
10	layer_64_3	sample_Resnet	0	5.669	1.306	0	0.227	5.954	0.06	0.03	0.008	0.03	0.06
11	layer_64_3	sample_Resnet	0	1.5945	0.08275	0	0.07075	0.5485	0.135	0.068	0.017	0.068	0.135
12	layer_64_3	sample_Resnet	0	0.383	0.409	0	0.227	1.108	0.06	0.03	0.008	0.03	0.06
13	layer_128_1	sample_Resnet	0	6.6015	0.4635	0	0.262	3.0585	0.477	0.239	0.06	0.239	0.477
14	layer_128_1	sample_Resnet	0	1.176	0.07225	0	0.06325	0.4915	0.12	0.06	0.015	0.06	0.12
15	layer_128_1	sample_Resnet	0	2.9015	0.31	0	0.2935	2.2045	0.537	0.269	0.068	0.269	0.537
16	layer_128_1	sample_Resnet	0	1.4375	0.14125	0	0.12925	0.9775	0.239	0.12	0.03	0.12	0.239
17	layer_128_2	sample_Resnet	0	1.3395	2.3275	0	0.905	4.1295	0.239	0.12	0.03	0.12	0.239
18	layer_128_2	sample_Resnet	0	2.9285	0.37	0	0.487	3.393	0.537	0.269	0.068	0.269	0.537
19	layer_128_2	sample_Resnet	0	6.5115	0.8545	0.8115	8.02825	20.68175	0.239	0.12	0.03	0.12	0.239
20	layer_128_3	sample_Resnet	0	4.6365	0.42475	0.01275	0.549	16.7855	0.239	0.12	0.03	0.12	0.239
21	layer_128_3	sample_Resnet	0	2.8865	1.8865	0	1.111	51.685	0.537	0.269	0.068	0.269	0.537

Overall Results


- Detailed.csv.

Average message latency per each collective phase (us)

	A	B	C	D	E	F	G	H	I	J	K	L	M
1			queuing d	queuing delay phase 1	queuing delay phase 2	queuing delay phase 3	queuing delay phase 4	queuing delay phase 5	network delay phase 1	network delay phase 2	network delay phase 3	network delay phase 4	network delay phase 5
2	conv1	sample_Resnet	0	1.8795	0.5455	0	0.137	2.3575	0.035	0.018	0.005	0.018	0.035
3	layer_64_1	sample_Resnet	0	0.453	0.03925	0	0.02725	0.2425	0.06	0.03	0.008	0.03	0.06
4	layer_64_1	sample_Resnet	0	2.9225	2.3655	0	0.065	7.1915	0.015	0.008	0.002	0.008	0.015
5	layer_64_1	sample_Resnet	0	6.8005	2.5125	1.2645	3.2625	6.7425	0.135	0.068	0.017	0.068	0.135
6	layer_64_1	sample_Resnet	0	2.013	0.412	0	0.227	2.536	0.06	0.03	0.008	0.03	0.06
7	layer_64_2	sample_Resnet	0	2.889	0.03925	0	0.02725	0.2425	0.06	0.03	0.008	0.03	0.06
8	layer_64_2	sample_Resnet	0	5.9785	0.08275	0	0.07075	0.5485	0.135	0.068	0.017	0.068	0.135
9	layer_64_2	sample_Resnet	0	1.901	2.947	0	0.227	4.313	0.06	0.03	0.008	0.03	0.06
10	layer_64_3	sample_Resnet	0	5.669	1.306	0	0.227	5.954	0.06	0.03	0.008	0.03	0.06
11	layer_64_3	sample_Resnet	0	1.5945	0.08275	0	0.07075	0.5485	0.135	0.068	0.017	0.068	0.135
12	layer_64_3	sample_Resnet	0	0.383	0.409	0	0.227	1.108	0.06	0.03	0.008	0.03	0.06
13	layer_128_1	sample_Resnet	0	6.6015	0.4635	0	0.262	3.0585	0.477	0.239	0.06	0.239	0.477
14	layer_128_1	sample_Resnet	0	1.176	0.07225	0	0.06325	0.4915	0.12	0.06	0.015	0.06	0.12
15	layer_128_1	sample_Resnet	0	2.9015	0.31	0	0.2935	2.2045	0.537	0.269	0.068	0.269	0.537
16	layer_128_1	sample_Resnet	0	1.4375	0.14125	0	0.12925	0.9775	0.239	0.12	0.03	0.12	0.239
17	layer_128_2	sample_Resnet	0	1.3395	2.3275	0	0.905	4.1295	0.239	0.12	0.03	0.12	0.239
18	layer_128_2	sample_Resnet	0	2.9285	0.37	0	0.487	3.393	0.537	0.269	0.068	0.269	0.537
19	layer_128_2	sample_Resnet	0	6.5115	0.8545	0.8115	8.02825	20.68175	0.239	0.12	0.03	0.12	0.239
20	layer_128_3	sample_Resnet	0	4.6365	0.42475	0.01275	0.549	16.7855	0.239	0.12	0.03	0.12	0.239


ResNet-50 Layer-Wise Raw Comm Latency

- A Torus 3D with total of 32 (2X4X4) nodes is used.
- Data parallel approach is used.
- Raw latency depends on the comm size plus the priority of each layer comm (queuing delay).




ResNet-50 Layer-Wise Compute vs. Exposed Comm Latency

- Exposed comm latency is observed for the first layer.
- because by the time we reach other layers except that.
first layer, their comm is already finished.



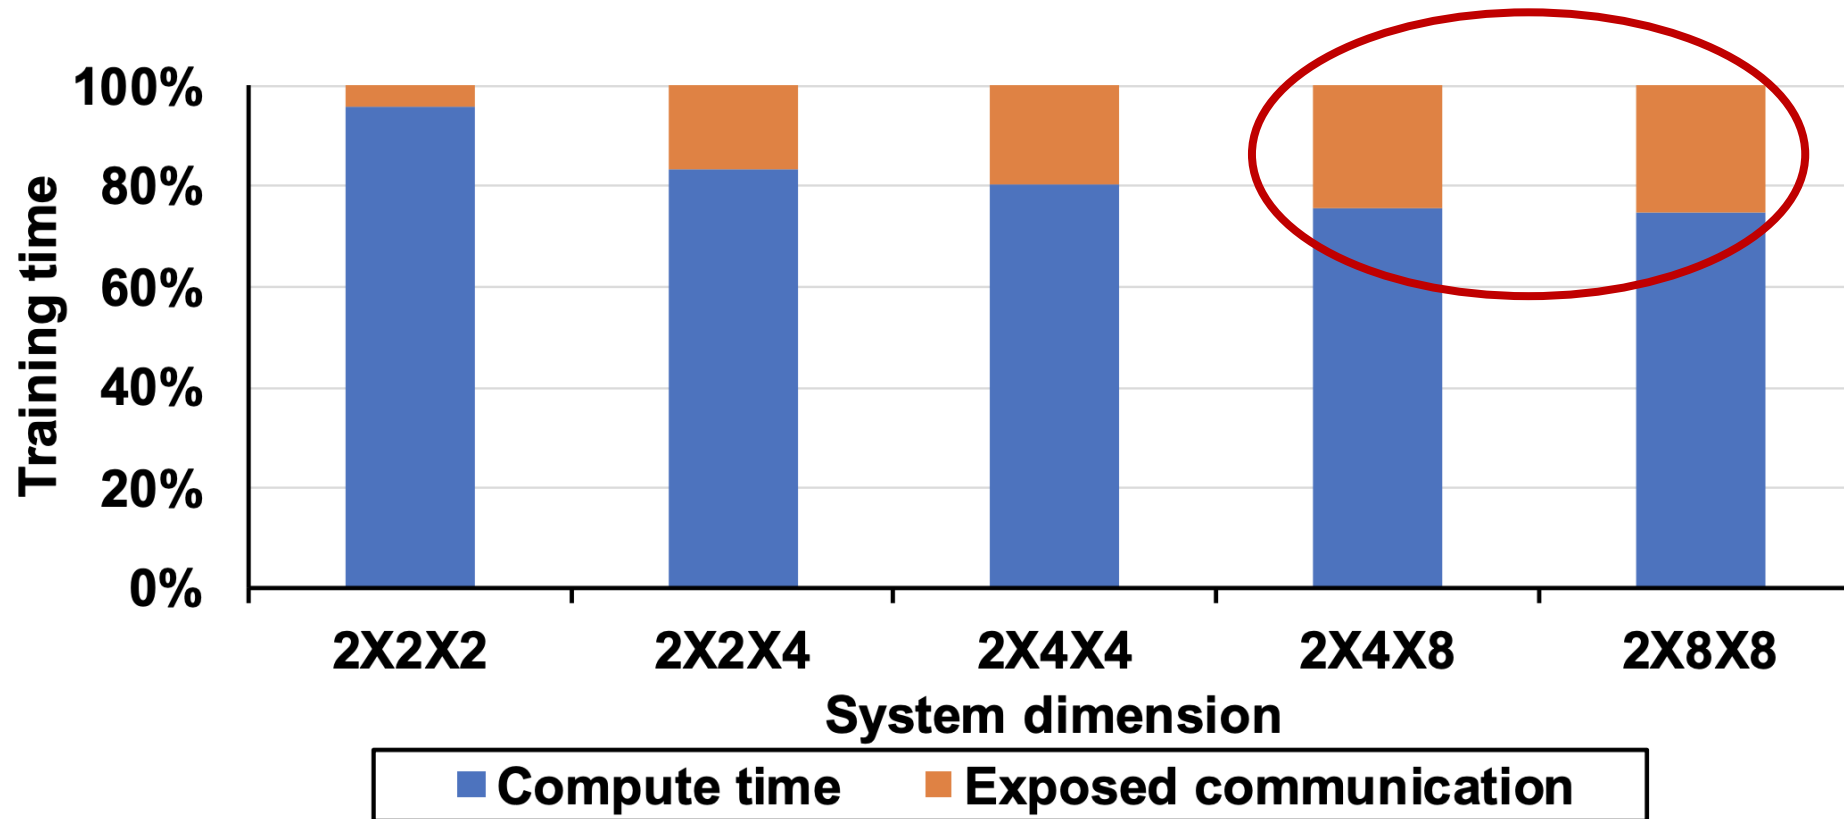
ResNet-50 Layer-Wise detailed latency

- Queue P2 is becoming the dominant factor due to very high speed of P1 (within package) that results most of the chunks get queued for the next phase (P2).




Effect of # of nodes on the Ratio of Total Compute vs Total Exposed Comm for ResNet-50

- A Torus 3D with total of 8, 16, 32, 64, 128 nodes are used.



Effect of Enhanced Compute Time per Node on the Ratio of Total Compute vs Total Exposed Comm for ResNet-50

- A Torus 3D with total of 32 nodes (2X4X4) is used.



Workload Generator

Workload Generator

- Generates workload inputs for the **text-based** workload engine
- Should receive the GEMM operations (M, N, K dimensions) and the parallelization strategy as input
- It uses SCALE-SIM simulator to find the compute times
- Please see [astra-sim/scripts/workload_generator/README.md](#)

Sample script to call workload generator

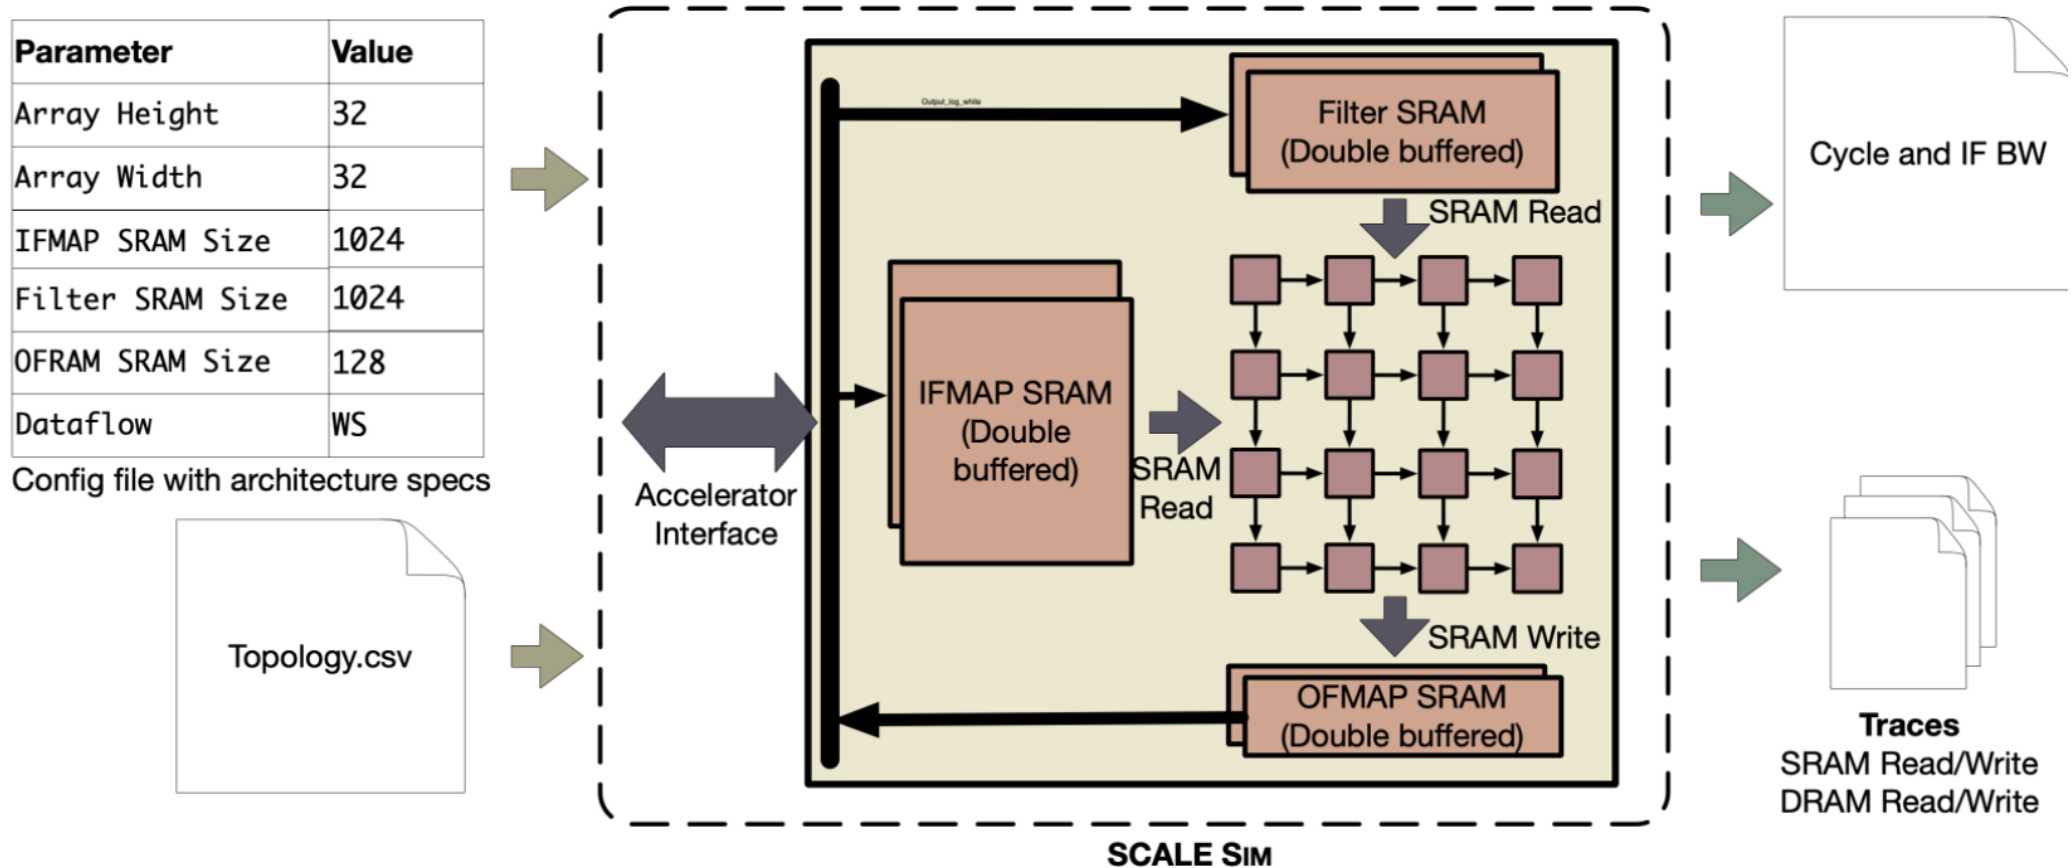
```
# For data-parallel
$ python3 gen_astrasim_workload_input.py \
  --datatype_size=2 \
  --mnk=mnk_inputs/example.csv \
  --num_npus=16 \
  --num_packages=2 \
  --output_file=../../inputs/workload/example_DATA.txt \
  --parallel=DATA \
  --run_name=example \
  --scalesim_config=../../extern/compute/SCALE-Sim/configs/google.cfg \
  --scalesim_path=../../extern/compute/SCALE-Sim
```

Sample MNK input file

	A	B	C	D	E
1	Layer	m	n	k	
2	MLP_Bottom	1024	128	512	
3	MLP_Bottom	1024	512	512	
4	MLP_Bottom	1024	512	512	
5	MLP_Bottom	1024	512	16	
6	MLP_Top_0	1024	1024	512	
7	MLP_Top_1	1024	512	512	
8	MLP_Top_2	1024	512	512	
9					
10					

SCALE-SIM

- <https://github.com/scalesim-project/scale-sim-v2>



Thank you!