


ASTRA-SIM Description

Network Layer



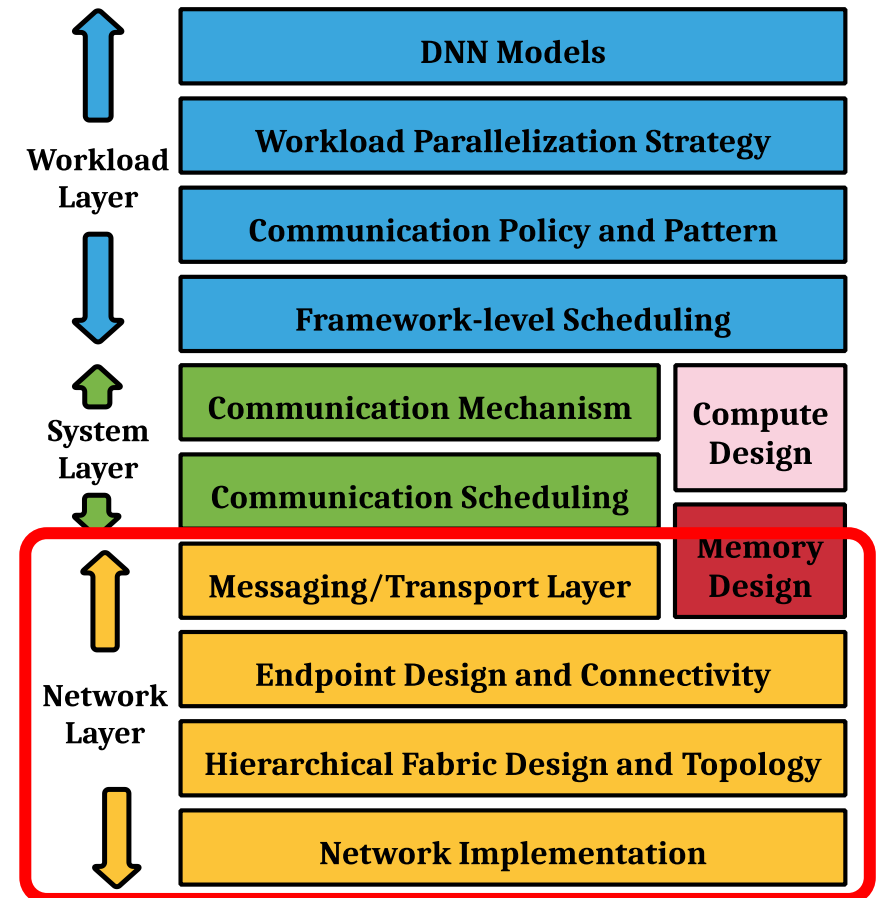
William Won

Ph.D. Student, School of Computer Science
Georgia Institute of Technology

Acknowledgments: Saeed Rashidi (Georgia Tech), Jinsun Yoo (Georgia Tech),
Srinivas Sridharan (Meta), Sudarshan Srinivasan (Intel)

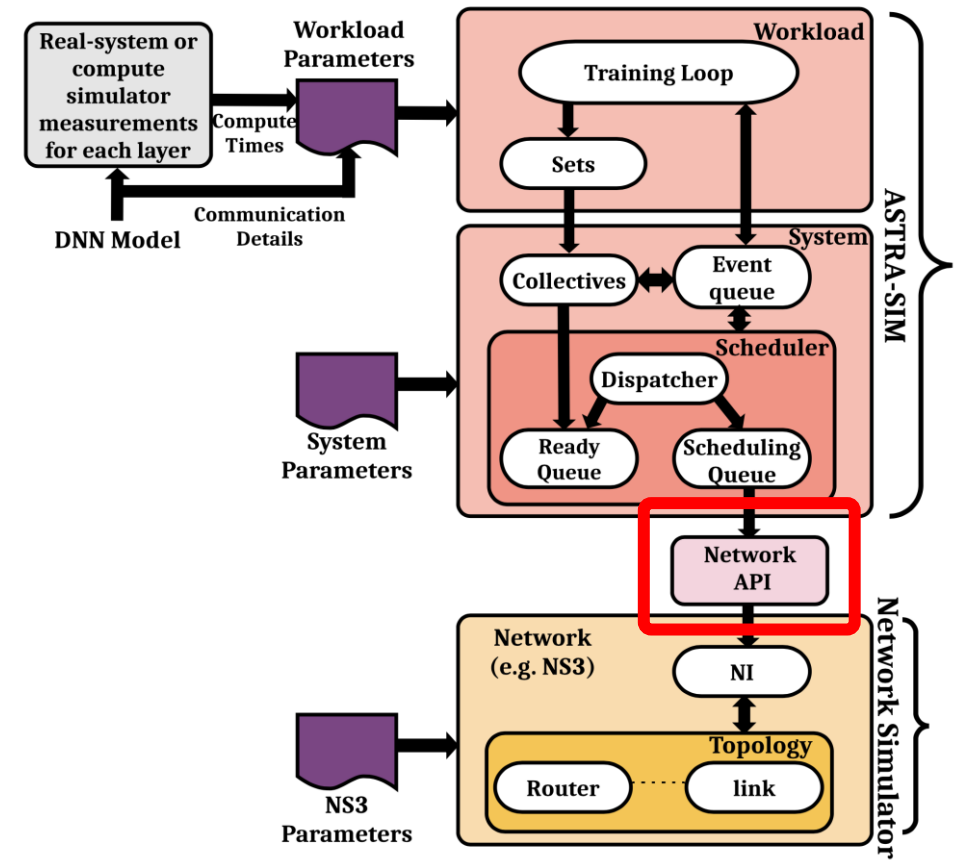
Network Layer

- Workload layer manages workload, parallelization, training loop
- System layer initiates collective communication
- Network layer simulates **actual network behaviors**
 - Communication protocols (TCP, RDMA, etc.)
 - Network topology
 - BW/latency per link
 - In-network collective communication
 - NIC offloading
 - Compression
 - Buffering, Arbitration



NetworkAPI

- Interface between System layer and Network backend
- Any network simulator implementing the NetworkAPI could be used as ASTRA-sim backend



(HOTI '20) Scalable Distributed Training of Recommendation Models: An ASTRA-SIM + NS3 case-study with TCP/IP transport

Example NetworkAPIs

- **sim_send(msg_size, src, dest, callback)**
 - Simulate sending a message of size msg_size from src through dest and **invoke callback function** once transmission has finished
- **sim_recv(msg_size, src, dest, callback)**
 - Simulate receiving a message of size msg_size from src through dest and **invoke callback function** once transmission has finished
- **sim_schedule(delta, callback)**
 - Invoke callback function after delta time
- **sim_get_time()**
 - Return current time of simulation to the frontend

Simulation Control Flow

- The main file (simulation entry point) is implemented inside network layer.
- Network layer creates corresponding **System** and **NetworkAPI** instances.
- Each system layer instance **internally** creates its workload layer instance.

Where do instantiations happen?

Analytical backend (analytical): analytical/analytical/main.cc

Analytical backend (congestion): analytical/congestion/main.cc

Garnet backend: garnet/gem5_astra/src/mem/ruby/network/garnet2.0/NetworkInterface.cc

ns3 backend: ns3-interface/simulation/scratch/AstraSimNetwork.cc

Network Layer

```
void main(){  
    .....  
    Instantiate NetworkAPI[];  
    Instantiate System[];  
    .....  
    for(auto &s:system){  
        s->workload.fire();  
    }  
    process_all_events();  
    return;  
}
```

Available Network Backends

- Network backends are maintained separately and are imported as **submodule**.
- We currently have **4 network backends** which implement NetworkAPI

Backend	Purpose	Notable Feature
analytical/analytical	analytical equation-based simulation	fast simulation, hierarchical topologies
analytical/congestion	congestion-aware analytical simulation	first-order congestion (queueing) modeling
Garnet	on-chip/scale-up network simulation	packetization, flow control, congestion
ns-3	inter-network simulation	large parallel GPU clusters

Analytical Backend

- Leverages **analytical equation** to estimate communication delay

- $$\text{delay}(\text{msg_size}, \text{src}, \text{dest}) = \underbrace{(\# \text{hops}) \times (\text{link latency})}_{\text{link delay}} + \underbrace{(\text{msg_size}) / (\text{link BW})}_{\text{serialization delay}}$$

- `sim_send(msg_size, src, dest, callback)`
 - Estimate communication delay
 - Assign callback to event queue after delay
- No congestion modeling
 - Appropriate for topology-aware collectives without network congestion
- **Fast simulation** for large-scale systems

(ISPASS '23) ASTRA-sim2.0: Modeling Hierarchical Networks and Disaggregated Systems for Large-model Training at Scale

Congestion-aware Analytical Backend

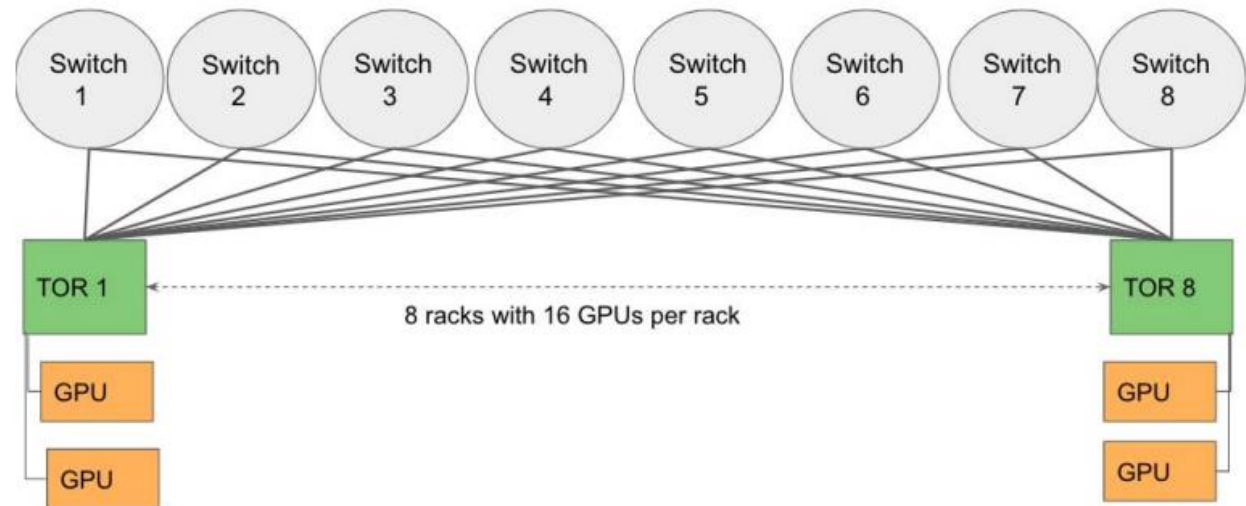
- First-order congestion modeling by per-link queueing
- Per-link delay is calculated using analytical equation
- e.g., `send(msg_size: 1 MB, route: [1, 2, 3, 4, 5])`
 - `send(1 MB, 1 → 2)`
 - `send(1 MB, 2 → 3)`
 - `send(1 MB, 3 → 4)`
 - `send(1 MB, 4 → 5)`
 - each send can be queued per each link
 - link processes pending chunks in-order
- **Fast simulation** for large-scale systems **with network congestion**

Garnet Backend

- Leverages **Garnet (interconnection network) simulator** as backend
- Appropriate for on-chip/scale-up networks
- Simulates interconnection network behaviors:
 - Message Packetization
 - Credit-based flow control
 - Congestion modeling
 - etc.
- Slower than analytical backend for large systems/models
- Supports switch-based/torus-based topologies

NS3 Backend

- Network simulator for **internet (inter-node) communication**
- Used to model ML training in **largely parallel GPU clusters**
- 1 NPU in Analytical/Garnet becomes **1 GPU** in NS3, connected with ToR/spine switch, etc.



(HOTI '22) Current RoCE congestion control methods have little impact on ML training workloads

Slide courtesy of: Jinsun Yoo (Georgia Tech)

Sample Network Input

Garnet network input

```
sample_torus.txt
1 num-npus: 12
2 num-packages: 6
3 package-rows: 3
4 topology: Torus3D
5 local-rings: 2
6 vertical-rings: 1
7 horizontal-rings: 1
8 flit-width: 2048
9 local-packet-size: 4096
10 package-packet-size: 4096
11 tile-link-width: 256
12 package-link-width: 256
13 vcs-per-vnet: 50
14 routing-algorithm: Ring_XY
15 router-latency: 1
16 local-link-latency: 90
17 package-link-latency: 200
18 buffers-per-vc: 5000
19 local-link-efficiency: 1.0
20 package-link-efficiency: 1.0
21
```

Analytical network input

```
sample_Torus3D.json
1 {
2   "topology-name": "Hierarchical",
3   "topologies-per-dim": ["Ring", "Ring", "Ring"],
4   "dimension-type": ["N", "N", "N"],
5   "dimensions-count": 3,
6   "units-count": [2, 2, 3],
7   "links-count": [2, 2, 2],
8   "link-latency": [10, 100, 100],
9   "link-bandwidth": [32, 16, 16],
10  "nic-latency": [0, 0, 0],
11  "router-latency": [0, 0, 0],
12  "hbm-latency": [500, 500, 500],
13  "hbm-bandwidth": [370, 370, 370],
14  "hbm-scale": [0, 0, 0]
15 }
```

